

OpenCL Host Framework

Easy Kernel Development, Zero Host Code

60-Minute Technical Presentation

Agenda

Part I: OCLEExample Introduction

- **Introduction / Features**
- **General OpenCL Host Flow**
- **Standard CL vs CL Extension**
- **Caching Mechanism**
- **JSON Configuration**

Part II: Two-Phase Architecture

- Recap
- Why Two-Phase Architecture
- Compilation Phase
- Execution Phase
- Cross-Platform Deliverable
- Summary & Q&A

The Problem We Solve

Traditional OpenCL Development Pain Points

- ✗ Writing 500+ lines of host code for each algorithm
- ✗ Managing OpenCL contexts, queues, buffers manually
- ✗ Compiling kernels and handling errors
- ✗ Setting up arguments and work sizes
- ✗ Verifying results against CPU reference
- ✗ Repeating this for every kernel variant

Solution: This framework eliminates all the boilerplate!

Introduction / Features

- The problem: 500+ lines of boilerplate for each OpenCL algorithm
- The solution: JSON-driven configuration, zero host code
- Six key benefits of the framework
- User journey: Build → Run → Learn
- High-level architecture overview

Framework Benefits



Zero Host Code

JSON-driven configuration



Instant Comparison

Run multiple variants



Platform Switching

Standard GPU vs CL Extension



Auto-Verification

GPU output vs CPU reference



Binary Caching

Fast re-runs with cached kernels



Easy Extension

Add new algorithms with scaffold

User Journey: Build → Run → Learn

User Journey: Build → Run → Learn



Available Algorithms (Each with Multiple Variants)

Resize

Image scale

v0: baseline
v1: local mem, v2: vectorized

SepFilter2D

Blur


v0: naive
v1: separable, v2: tiled

FullyConnected

Matrix multiply

v0: naive
v1: tiled, v2: vectorized

High-Level Architecture




General OpenCL Host Flow

- The 6-step standard OpenCL host flow
- What the framework handles automatically
- Key timing: ~100ms compilation overhead
- Why this matters for iteration speed

Standard OpenCL Host Flow

Standard OpenCL Host Flow (6 Steps)



Framework Components

algorithm_runner.c
Main orchestrator

opencl_utils.c
OpenCL abstraction

kernel_args.c
Argument parsing

cache_manager.c
Binary caching

config.c
JSON config parser

verification.c
Result verification

op_registry.c
Algorithm registry

All components work together seamlessly - you just run: `./run.sh <algo> <variant>`


Standard CL vs CL Extension

- Two platform types: Standard OpenCL and CL Extension
- Standard CL: Works on any GPU, great for development
- CL Extension: Vendor-optimized, maximum performance
- One-line config change to switch between them

Platform Branching

Standard OpenCL vs CL Extension

Shared flow with branching only at kernel execution



← Steps 1-4 → Decision → Step 5a/5b → Step 6 →

Platform Configuration in JSON

```
"kernels": {  
    "v0": {  
        "host_type": "standard",  
        "kernel_file": "resize1.cl"  
    },  
    "v1": {  
        "host_type": "cl_extension",  
        "kernel_file": "resize2.cl"  
    }  
}
```

Key Benefits

- ✓ Runtime platform selection
- ✓ No code changes required
- ✓ No rebuild needed
- ✓ Switch with one line change
- ✓ Run both variants to compare
- ✓ Same kernel.cl works on both


Easy to switch between Standard CL and CL Extension

Caching Mechanism

- Binary caching eliminates repeated compilation
- Cache miss: First run compiles (~100ms), saves binary
- Cache hit: Load cached binary (~5ms) — 20x faster!
- Automatic invalidation when source changes (SHA256)

Binary Caching

OpenCL Kernel Binary Cache Mechanism



Performance Gain

First run: ~100ms (compile)

Cached: ~5ms (load binary)

20x faster!

How It Works

1. Hash source code to detect changes
2. If hash matches -> load cached binary
3. If hash differs -> recompile and update cache

Code Reference

cache_manager.c
opencl_utils.c:438-519
Hash: FNV-1a (16 bytes)

JSON Configuration

- Configuration file structure in config/ directory
- Global inputs.json and algorithm-specific configs
- Complete kernel configuration with all parameters
- Six supported argument types for any kernel pattern

```
config/
├── inputs.json          # Global input definitions
├── resize.json          # Algorithm-specific
├── sepfilter2d.json     # Algorithm-specific
└── fullyconnected.json  # Algorithm-specific
```

Kernel Configuration

```
"kernel": {  
    "description": "standard cl implementation",  
    "host_type": "standard",  
    "kernel_option": "-g",  
    "kernel_file": "examples/resize/cl/resize0.cl",  
    "kernel_function": "resize",  
    "work_dim": 2,  
    "global_work_size": [1920, 1088],  
    "local_work_size": [16, 16],  
    "kernel_args": [  
        {"i_buffer": ["uchar", "src"]},  
        {"o_buffer": ["uchar", "dst"]},  
        {"param": ["int", "src_width"]}  
    ]  
}
```

Supported Argument Types

| Type | Syntax | Description |
|---------------|-----------------------------------|-------------------|
| Input Buffer | {"i_buffer": ["uchar", "src"]} | Read-only image |
| Output Buffer | {"o_buffer": ["uchar", "dst"]} | Write-only result |
| Parameter | {"param": ["int", "width"]} | Auto from image |
| Scalar | {"scalar": ["float", "sigma"]} | Custom value |
| Custom Buffer | {"buffer": ["float", "wts", 100]} | User-defined |
| Struct | {"struct": ["a", "b", "c"]} | Packed scalars |

Part I Summary

| Topic | Key Points |
|-------------------------|---|
| Introduction / Features | Zero host code, JSON-driven configuration |
| OpenCL Host Flow | 6-step flow handled by framework |
| Platform Options | Standard CL vs CL Extension — one-line switch |
| Caching | Binary caching: 100ms → 5ms |
| Configuration | JSON-based kernel and I/O setup |

Any questions before Part II?

Agenda

Part I: OCLEExample Introduction

- Introduction / Features
- General OpenCL Host Flow
- Standard CL vs CL Extension
- Caching Mechanism
- JSON Configuration

Part II: Two-Phase Architecture

- **Recap**
- **Why Two-Phase Architecture**
- **Compilation Phase**
- **Execution Phase**
- **Cross-Platform Deliverable**
- **Summary & Q&A**

Part I Recap

Framework Benefits

Zero host code, JSON-driven

User Journey

Build → Run → Learn

Platform Options

Standard CL vs CL Extension

JSON Configuration

Kernel config, I/O setup

Now: How we split the host into two phases


Why Two-Phase Architecture

- ✗ Every run recompiles kernel (~100ms overhead)
- ✗ Source code must be present on target
- ✗ Can't pre-compile for different devices
- ✗ No binary distribution possible

Solution: Split into Compilation + Execution phases

Two-Phase Architecture

Two-Phase Architecture: Compile Once, Run Anywhere



Compilation Phase (WIP)

- Runs once on Linux development machine
- Input: kernel.cl source + config.json
- Output: kernel.bin (compiled) + gpu_host (executor) OR mvpu_host (executor)
- kernel.bin includes metadata for target device (how about batch binary ?!)

Compilation Phase (WIP)

Purpose: Convert source + config into deliverable-ready binaries

Inputs

- kernel.cl — Source code
- config.json — Configuration

Outputs

- kernel.bin — Compiled binary
- gpu_host — Linux executable
- mvpu_host (Player) — Linux executable
- mvpu_host (Player) — Android executable

kernel.bin contains:

- Compiled kernel binary (device-specific)
- Metadata (device name, OpenCL version)
- Compile options used

Execution Phase (WIP)

- Runs many times on target device (Linux/Android)
- Input: kernel.bin (pre-compiled) + config + data
- Uses `clCreateProgramWithBinary()` — no compilation!
- Performance: ~5ms vs ~105ms — 20x faster startup

Execution Phase (WIP)

Purpose: Run kernel without compilation overhead

Inputs

- kernel.bin — Pre-compiled
- config.json — Runtime config
- Input data

Output

- result.bin — Processed output
- Using `clCreateProgramWithBinary()`

Performance: ~5ms vs ~105ms (with compilation)


Loading pre-compiled binary eliminates 100ms compilation overhead

Cross-Platform Deliverable

- Develop and compile on Linux (full toolchain)
- Deploy to Linux (GPU/MVPU) OR Android (MVPU) devices
- Only 3 types of files: `opencl_host(s)` + `kernel.bin` + `config.json`
- Deliver offline compiler: No host source code, headers are needed on target

Cross-Platform Deliverable

Cross-Platform Deployment



What to Deliver

Execution Package:

```
out/
└── opencl_host(s)      # Executor binary
└── kernel.bin          # Compiled kernel
└── config.json          # Runtime config
```

NOT required in execution:

- ✗ kernel.cl (source code)
- ✗ include/ (headers)
- ✗ Host code build tools
- ✗ Host code development dependencies

Benefits: Smaller footprint • No source exposure • Faster startup

Summary

| Phase | When | Where | Output |
|-------------|------------|---------------|--------------------------|
| Compilation | Once | Linux | kernel.bin + opencl_host |
| Execution | Many times | Linux/Android | result.bin |

Key Benefits:

- ✓ Zero compile time after first run
- ✓ Binary-only deliverable
- ✓ Cross-platform support (GPU + MVPU)
- ✓ Smaller execution footprint

Questions & Discussion

Thank you!